

# Mapping Unikernels with TAG based architectures



**Akilan Selvacoumar**

Mathematics and Computer Sciences  
Heriot Watt University

Year 1 progression report of:  
*Doctor of Philosophy*

October 2022



I would like to dedicate this thesis to my loving parents ...



## **Declaration**

I hereby declare that except where specific reference is made to the work of others, the contents of this dissertation are original and have not been submitted in whole or in part for consideration for any other degree or qualification in this, or any other university. This dissertation is my own work and contains nothing which is the outcome of work done in collaboration with others, except as specified in the text and Acknowledgements. This dissertation contains fewer than 65,000 words including appendices, bibliography, footnotes, tables and equations and has fewer than 150 figures.

Akilan Selvacoumar

October 2022



## **Acknowledgements**

And I would like to acknowledge ...





## **Abstract**

This is where you write your abstract ...



# Table of contents

<b>List of figures</b>	<b>xv</b>
<b>List of tables</b>	<b>xvii</b>
<b>Nomenclature</b>	<b>xix</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 Motivation</b>	<b>3</b>
<b>3 Research Questions</b>	<b>5</b>
<b>4 Literature Review</b>	<b>7</b>
4.1 TAG based architecture survey . . . . .	7
4.1.1 Timder V . . . . .	7
4.1.2 ARM MTE . . . . .	8
4.1.3 D-RI5CY . . . . .	8
4.1.4 TMDFI . . . . .	8
4.1.5 HyperFlow . . . . .	8
4.1.6 SDMP . . . . .	9
4.1.7 Typed Architecture . . . . .	10
4.1.8 Dover . . . . .	12
4.1.9 Shakti-T . . . . .	12
4.1.10 HDFI . . . . .	12
4.1.11 lowRISC . . . . .	12
4.1.12 Taxi . . . . .	12
4.1.13 Pump . . . . .	12
4.1.14 CHERI . . . . .	12
4.1.15 SPARC M7/M8 SSM . . . . .	12

4.1.16	Low-Fat Pointers . . . . .	12
4.1.17	SAFE . . . . .	12
4.1.18	DataSafe . . . . .	12
4.1.19	Harmoni . . . . .	12
4.1.20	Shioya, et al. . . . .	12
4.1.21	SIFT . . . . .	12
4.1.22	FlexCore . . . . .	12
4.1.23	Execution Leases . . . . .	12
4.1.24	GLIFT . . . . .	12
4.1.25	TIARA . . . . .	12
4.1.26	DIFT Coprocessor . . . . .	12
4.1.27	HardBound . . . . .	12
4.1.28	Loki . . . . .	12
4.1.29	FLexiTaint . . . . .	12
4.1.30	SECTAG . . . . .	12
4.1.31	Raksha . . . . .	12
4.1.32	SecureBit . . . . .	12
4.1.33	Minos . . . . .	12
4.1.34	DIFT . . . . .	12
4.1.35	RIFLE . . . . .	12
4.1.36	AEGIS . . . . .	12
4.1.37	Mondriaan . . . . .	12
4.1.38	Aries . . . . .	12
4.1.39	XOM . . . . .	12
4.1.40	M-Machine . . . . .	12
4.1.41	KCM . . . . .	12
4.1.42	SPUR . . . . .	12
4.1.43	Lisp Machine . . . . .	12
4.1.44	HEP . . . . .	12
4.1.45	Burroughs . . . . .	12
<b>5</b>	<b>Expirements</b>	<b>13</b>
<b>6</b>	<b>Research Goals</b>	<b>15</b>
<b>7</b>	<b>Research Timeline</b>	<b>17</b>

Table of contents	xiii
<b>8 Conclusion</b>	<b>19</b>
<b>References</b>	<b>21</b>
<b>References</b>	<b>23</b>
<b>Index</b>	<b>25</b>



## List of figures





## List of tables



# Nomenclature

## Roman Symbols

$F$  complex function

## Greek Symbols

$\gamma$  a simply closed curve on a complex plane

$\iota$  unit imaginary number  $\sqrt{-1}$

$\pi$   $\simeq 3.14\dots$

## Superscripts

$j$  superscript index

## Subscripts

0 subscript index

crit Critical state

## Other Symbols

$\oint_{\gamma}$  integration around a curve  $\gamma$

## Acronyms / Abbreviations

ALU Arithmetic Logic Unit

BEM Boundary Element Method

CD Contact Dynamics

CFD Computational Fluid Dynamics

<i>CIF</i>	Cauchy's Integral Formula
CK	Carman - Kozeny
DEM	Discrete Element Method
DKT	Draft Kiss Tumble
DNS	Direct Numerical Simulation
EFG	Element-Free Galerkin
FEM	Finite Element Method
FLOP	Floating Point Operations
FPU	Floating Point Unit
FVM	Finite Volume Method
GPU	Graphics Processing Unit
LBM	Lattice Boltzmann Method
LES	Large Eddy Simulation
MPM	Material Point Method
MRT	Multi-Relaxation Time
PCI	Peripheral Component Interconnect
PFEM	Particle Finite Element Method
PIC	Particle-in-cell
PPC	Particles per cell
RVE	Representative Elemental Volume
SH	Savage Hutter
SM	Streaming Multiprocessors
USF	Update Stress First
USL	Update Stress Last

# **Chapter 1**

## **Introduction**



# **Chapter 2**

## **Motivation**





## **Chapter 3**

### **Research Questions**



# Chapter 4

## Literature Review

The literature review is split into 3 sections. The first section talks about the papers surveyed for Unikernels and the 2nd section talks about papers surveyed for TAG based architectures and the third sections talks about the possible incentives of combining them both which helps answer the research questions stated (TODO: Add reference to research question section).

### 4.1 TAG based architecture survey

The following was a survey conducted on exisisting TAG based implementations and the recent survey based on TAG based architectures (//TODO add survey reference) published in 2022 was a good staring point to understand about various implementations of TAG based architectures with the high level metrits and limitations. The following section provides our own version of the Survey to help decide the best implementations to answer the research questions (//TODO reference research questions chapter).

According to the TAG based architecture survey (//TODO add survey reference) there are 37 published efforts on TAG based architectures over the past decade and 20 published efforts preceding that.

#### 4.1.1 Timder V

It is a tagged memory architecture for flexible and efficient isolation of code and data on small embedded systems. The TAG isolation is augmented with a memory protection unit to isolate induvidual processes. Timber V is compatible with exsisting code. The contributions of the paper are:

- Efficient tagged memory architecture for isolated execution on low-end processors.

- Concept introduced called stack interleaving that allows efficient and dynamic memory management.
- Lightweight shared memory between enclaves.
- Efficient shared MPU (i.e Memory Protection Unit) design.

#### **4.1.2 ARM MTE**

The ARMv8.5-Memory Tagging Extension (MTE) aims to increase the memory safety written for unsafe languages without requiring source code changes and in certain cases without recompilation. It generally focuses on the bounds checking use case, Though it provides limited tags which means it can only provide probabilistic overflow detection. It is one of the latest commercial incarnations of memory-safety-focused tagged architectures.

#### **4.1.3 D-RI5CY**

It provides a design and implementation of a hardware dynamic information flow tracking (DIFT) architecture for RISC-V processor cores. The paper presents a low overhead implementation of DIFT that is specialized for low-end embedded systems for IOT applications. The following are high level contributions:

- Design of D-RI5CY, A DIFT-protected implementation of the RI5CY processor core. The paper implements the modification of the DIFT TAG propagation and TAG checking mechanism in a way that is transparent to the execution of the regular instructions.
- Concept introduced called stack interleaving that allows efficient and dynamic memory management.
- Lightweight shared memory between enclaves.
- Efficient shared MPU (i.e Memory Protection Unit) design.

#### **4.1.4 TMDFI**

#### **4.1.5 HyperFlow**

It is a design and security implementation that offers security assurance because it is implemented using a security-typed hardware description language. It allows complex information

flow policies to be configured at run time. The paper introduces ChiselFlow, a new secure hardware description language. The contribution of the paper includes:

- Processor architecture and implementation designed for timing-safe information flow security.
- Complete RISC-V instruction set extended with instructions for information flow control.
- Verified at design time with a hardware description language.
- Novel representations of lattices that can be implemented in hardware efficiently.

Hyperlow implements a nonmalleable IFC policy using tags. To eliminate timing side channels, the processor tracks the tag of the currently executing code and flushes caches, TLB, branch predictor, and other microarchitectural state on changes in the confidentiality or integrity tag of the running code. The modifications to avoid timing side channels seem more extensive than those to add tags. The authors report overheads in cycles per instruction of between 1 operation to the worst-case number of cycles.

#### 4.1.6 SDMP

This paper focuses on designing metadata tag based stack-protection security policies for general purpose tagged architecture. The policies specifically exploit the natural locality of dynamic program call graphs to achieve cacheability of the metadata rules that they require. The simple Return Address Protection policy has a performance overhead of 1.2% but just protects return addresses. The two richer policies present, Static Authorities and Depth Isolation, provide object-level protection for all stack objects. When enforcing memory safety, The Static Authorities policy has a performance overhead of 5.7% and our Depth Isolation policy has a performance overhead of 4.5%. The contribution of the paper includes:

- The formulation of a range of stack protection policies within the SDMP model.
- Three optimizations for our stack policies: Lazy Tagging, Lazy Clearing and Cache Line Tagging.
- The performance modeling results of our policies on a standard benchmark set, including the impact of our proposed optimizations.

### 4.1.7 Typed Architecture

This paper introduces Typed Architectures, a high-efficiency, low-cost execution substrate for dynamic scripting languages, where each data variable retains high-level type information at an ISA level. Typed Architectures calculate and check the dynamic type of each variable implicitly in hardware, rather than explicitly in software. Typed Architectures provide hardware support for flexible yet efficient type tag extraction and insertion, capturing common data layout patterns of tag-value pairs. The evaluation using a fully synthesizable RISC-V RTL design on FPGA shows that Typed Architectures achieve mean speedups of 11.2% and 9.9% with minimum speedups of 32.6% for JavaScript and Lua. The contribution of the paper includes:

- ISA extension to efficiently manage type tags in hardware, which can be flexibly applied to multiple scripting languages and engines.
- Design and implement the Typed Architecture pipeline, which effectively reduces the overhead of dynamic type checking at low hardware cost.
- Prototype the proposed processor architecture using a fully synthesizable RTL model to execute two production-grade scripting engines with large inputs on FPGA (executing over 274 billion instructions in total) and provide a more accurate estimate of area and power using a TSMC 40nm standard cell library.



**4.1.8 Dover**

**4.1.9 Shakti-T**

**4.1.10 HDFI**

**4.1.11 lowRISC**

**4.1.12 Taxi**

**4.1.13 Pump**

**4.1.14 CHERI**

**4.1.15 SPARC M7/M8 SSM**

**4.1.16 Low-Fat Pointers**

**4.1.17 SAFE**

**4.1.18 DataSafe**

**4.1.19 Harmoni**

**4.1.20 Shioya, et al.**

**4.1.21 SIFT**

**4.1.22 FlexCore**

**4.1.23 Execution Leases**

**4.1.24 GLIFT**

**4.1.25 TIARA**

**4.1.26 DIFT Coprocessor**

**4.1.27 HardBound**

**4.1.28 Loki**

**4.1.29 FLexiTaint**

**4.1.30 SECTAG**

**4.1.31 Raksha**

**4.1.32 SecureBit**



# **Chapter 5**

## **Experiments**



## **Chapter 6**

### **Research Goals**



## **Chapter 7**

### **Research Timeline**



## **Chapter 8**

## **Conclusion**





## References



## References



# **Index**

LaTeX class file, 1